

Identifying Tumor Growth Drivers in Biomedical Text

Nathan Byers

10/26/2017

Motivation

- ▶ Did things backwards: picked my project first, then looked for paper. . .
- ▶ Personalized medicine Kaggle competition
- ▶ “Classifying Clinically Actionable Genetic Mutations”
- ▶ <https://www.kaggle.com/c/msk-redefining-cancer-treatment>

Problem

- ▶ A sequenced tumor could have thousands of mutations
- ▶ Which mutations drive tumor growth and which don't

driver <http://www.pancreaticcancer.net.au/research-genomics/>

Problem

- ▶ How is this currently handled in a precision medicine context?
- ▶ Manual process of searching text in biomedical literature to classify mutations
- ▶ This is a bottleneck
- ▶ Can it be automated?

Kaggle Challenge

- ▶ Come up with an algorithm that will automate this process
- ▶ A training data set of articles where the tumors have been manually classified is provided on the Kaggle website
- ▶ A test data set is provided to score your algorithm
- ▶ My project will be to develop an algorithm that will classify the mutations in the Kaggle training data and evaluate the performance on the test data

Paper

- ▶ Singhal A, Simmons M, and Lu Z. “Text mining for precision medicine: automating disease-mutation relationship extraction from biomedical literature.” *Journal of the American Medical Informatics Association* (2016) 23: 766-772.
- ▶ Objective
 - ▶ Identify disease-mutation relationships from biomedical texts
 - ▶ Develop a machine learning algorithm to automate the process
 - ▶ Focused on breast and prostate cancers

Materials

- ▶ Manually annotated text for breast and prostate cancers from Doughty *et al.* for training the model
- ▶ Test set built from PubMed
- ▶ tmVar tool used to identify mutations in the text (Wei *et al.*)
- ▶ DNorm tool used to identify the disease names in the text (Leaman *et al.*)
- ▶ Weka for prediction

methods1

methods2

Feature Construction

- ▶ Target disease frequency score (TDFS)
- ▶ “This score is computed as the frequency count for the target disease mentioned in the input text. This feature adds information about the dominance of target disease mentions in the text.”
- ▶ Nearness to target disease score (NTDS)
- ▶ “For a mutation identified in the text, its NTDS is an integer denoting a cumulative score of all the times this mutation has the target disease as the closest disease mentioned in the text.”

Feature Creation

- ▶ Other disease frequency score (ODFS)
- ▶ “Unlike the TDFS, which captures information about the target disease in the text, the ODFS denotes the frequency of the next most frequent disease mention in the text other than the target disease.”
- ▶ Same sentence disease-mutation co-occurrence score (DMCS)
- ▶ “For a mutation name and its nearest disease mentioned in the text, the DMCS is a binary score denoting the co-occurrence of the mutation and its nearest disease in the same sentence.”

Feature Creation

- ▶ Within text sentiment score
- ▶ “For a mutation name in the text and its corresponding nearest disease mentioned in the text, we extracted the ‘within text,’ which refers to the text between the mutation and the nearest disease mentioned. This text is then analyzed for its sentiment. . .”
- ▶ Text sentiment subjectivity score (TSSS)
- ▶ “The TSSS corresponds to the subjectivity of the sentiment score computed in the previous feature. It provides an estimate of the reliability of the sentiment score.”

ML Algorithm

- ▶ Tested decision tree, multilayer perceptron, and Bayesian logistic regression
- ▶ Only reported decision tree results because of its superior performance

Results: Prostate Cancer

	EMU*	tmVar+ML
Precision	0.729	0.904
Recall	0.803	0.856
F-measure	0.764	0.880

- ▶ Current state of the art mutation extraction for given disease (Doughty *et al.*)

Results: Breast Cancer

	EMU	tmVar+ML
Precision	0.806	0.878
Recall	0.852	0.813
F-measure	0.828	0.845

My Project

- ▶ Identify diseases (DNorm) and mutations (tmVar) in the Kaggle training text using the National Center for Biotechnology Information API
- ▶ <https://www.ncbi.nlm.nih.gov/research/bionlp/APIs/>
- ▶ Create features
- ▶ Frequency of mutation
- ▶ Distance from mutation to disease name
- ▶ Distance from other mutations
- ▶ Sentiment of text

My Project

- ▶ Apply several machine learning algorithms
- ▶ Decision Tree
- ▶ Neural Network
- ▶ Support Vector Machine
- ▶ Score on the test text and see which is most accurate

References

- ▶ Doughty E, Kertesz-Farkas A, Bodenreider O, et al. "Toward an automatic method for extracting cancer- and other disease-related point mutations from the biomedical literature." *Bioinformatics*. 2011 27(3):408–415
- ▶ Singhal A, Simmons M, and Lu Z. "Text mining for precision medicine: automating disease-mutation relationship extraction from biomedical literature." *Journal of the American Medical Informatics Association* 2016 23: 766-772
- ▶ Leaman R, Doğan RI, Lu Z. "DNorm: disease name normalization with pairwise learning to rank." *Bioinformatics* 2013 29(22):2909–2917
- ▶ Wei C-H, Harris BR, Kao H-Y, Lu Z. "tmVar: a text mining approach for extracting sequence variants in biomedical literature." *Bioinformatics* 2013 29(11):1433–1439